

# Secure Two Party Computation

## A practical comparison of recent protocols

Nick Tutte

University of Bristol

---

# Project Background

- Secure Multi Party Computation (SMC) is a long standing problem in cryptography.
  - ▶ How can a set of mutually distrusting parties collaborate to compute a function without revealing their inputs?
- We focused on the sub-problem of Secure *Two* Party Computation.
- Yao Garbled Circuits are a popular approach to S2C, but naive Yao Circuits only provide security against passive adversaries.
- Recently several protocols have been proposed to provide security against active adversaries.
- Our aim was to provide the first practical comparison of some of these protocols performance.

---

# Project Motivation

- ✿ Without a good comparison of these protocols we cannot know which is the most promising for further research.
- ✿ Whilst limited theoretical comparisons have been made these can only take us so far.
- ✿ In order to run practical comparisons we need implementations to compare.
- ✿ Furthermore we broke down our results to identify the bottlenecks for each protocol.

---

# Secure Two Party Computation

- ✿ Two parties wish to compute a function  $(y_1, y_2) = f(x_1, x_2)$
- ✿ They desire the following security properties,
  - ▶ Privacy - Parties only learn their output from participating.
  - ▶ Correctness - Output is correct evaluation of the function *given the inputs*.
  - ▶ Independence of inputs - Neither party can give their input as a function of the others.

---

# Yao's Protocol

- ✿ Proposed by Andrew Yao in 1986. Primarily used for Two Party computation but it can be extended to Multiparty computation.
- ✿ The basic idea,
  - ▶ One party is labelled the 'Builder' and the other is the 'Executor'.
  - ▶ The Builder constructs a binary circuit representing the function.
  - ▶ The Builder then 'garbles' this circuit, encrypting it in such a way that it can be evaluated without revealing information about the inputs.

---

# The Security of Yao's Protocol

- ✿ As already mentioned Yao's protocol is only secure against passive adversaries.<sup>1</sup>
- ✿ The Executor has no way to know if the circuit given to it by the Builder actually computes the function it purports to.
- ✿ There are two main approaches to overcoming this,
  - ▶ Cut and Choose.
  - ▶ Commit and Prove.

---

<sup>1</sup>For a formal proof of this see "A Proof of Security of Yao's Protocol for Two-Party Computation" by Lindell and Pinkas (2006).

---

## Cut and Choose - Concept

- ✿ All three protocols we implemented are based on Cut and Choose.
- ✿ The Builder now generates many garbled circuits and sends them to the Executor.
- ✿ The Executor then picks a subset of these circuits and asks the Builder to open them so they can be checked for correctness.
- ✿ If all check circuits pass then the remaining circuits are evaluated.

---

## Cut and Choose - Pitfalls

- ✶ Cut and Choose seems trivially simple, however it actually raises several new problems. The main ones are,
  - ▶ Consistency of inputs - Ensuring the parties give the same input to the many circuits, without leaking the inputs to the other party.
  - ▶ Output determination - Now we have many circuits what do we do if their outputs disagree?
- ✶ Protocols aiming to provide active security must address each of these issues.





---

## Lindell-Pinkas 2010 (LP-2010)

- ✿ Based around a new primitive, a Cut and Choose Oblivious Transfer.
- ✿ Using an elegant approach to consistency using Zero Knowledge proofs.
- ✿ Half of the circuits are check circuits, the other half are evaluation circuits.
- ✿ Return majority output of the evaluation circuits.
- ✿ Statistical security  $2^{-0.311 \cdot S}$  where  $S$  is the number of circuits.
- ✿ So we need 130 circuits to achieve statistical security of  $2^{-40}$ .

---

## Lindell 2013 (L-2013)

- ✿ Builds on Lindell-Pinkas 2010 using the same basic ideas.
- ✿ Modifies the Cut and Choose Oblivious Transfer.
- ✿ Runs a small sub-computation (using Lindell-Pinkas) such that only one circuit needs to be correct.
- ✿ We expected it to perform well for large circuits and poorly for small circuits.
- ✿ Statistical security  $2^{-S}$ , so we need 40 circuits to achieve statistical security of  $2^{-40}$ .

---

## Huang-Katz-Evans 2013 (HKE)

- Developed concurrently to Lindell 2013.
- Uses Symmetric cut and choose. Both parties build circuits, both parties execute the other's circuits.
- Uses a logarithm based approach to the consistency of inputs.
- Output determination is such that,
  - For each output wire a value is output if at least one of your circuits gives that value *and* at least one of you partner's circuits gives the same value.
- Statistical security  $2^{-S+\log(S)}$ , so *each* party needs 46 circuits to achieve statistical security of  $2^{-40}$ .

---

## Merging L-2013 and HKE (L-HKE)

- ✿ The natural question raised by L-2013 is can we improve it by changing the protocol used for the sub-computation?
- ✿ We changed the sub-computation to use HKE instead.
- ✿ Not a trivial task,
  - ▶ Lindell provides several levels of optimisations to the sub-computation some of which are incompatible with HKE.
  - ▶ The Lindell-Pinkas Zero Knowledge Proof approach to cannot work here, we have to switch the whole protocol to use the HKE logarithm based approach.
  - ▶ Additionally the output of the sub-computation needs to be hidden from the Builder.

## 32-bit Addition Circuit (439 gates)

| Builder | CPU Time | Wall Time | Bytes Sent | Bytes Recv |
|---------|----------|-----------|------------|------------|
| LP-2010 | 113.96   | 27.41     | 7,648,074  | 737,109    |
| L2013   | 171.21   | 42.03     | 4,693,761  | 980,193    |
| HKE     | 45.59    | 6.77      | 3,143,383  | 3,143,366  |
| L-HKE   | 145.77   | 25.47     | 5,995,366  | 3,299,399  |

| Executor | CPU Time | Wall Time | Bytes Sent | Bytes Recv |
|----------|----------|-----------|------------|------------|
| LP-2010  | 55.90    | 27.45     | 737,109    | 7,648,074  |
| L-2013   | 101.69   | 42.05     | 980,193    | 4,693,761  |
| L-HKE    | 132.51   | 25.83     | 3,299,399  | 5,995,366  |

HKE runs fastest by a very clear margin, while L-2013 uses the least bandwidth.

## 32-bit Addition Circuit (439 gates)

| Builder | CPU Time | Wall Time | Bytes Sent | Bytes Recv |
|---------|----------|-----------|------------|------------|
| LP-2010 | 113.96   | 27.41     | 7,648,074  | 737,109    |
| L2013   | 171.21   | 42.03     | 4,693,761  | 980,193    |
| HKE     | 45.59    | 6.77      | 3,143,383  | 3,143,366  |
| L-HKE   | 145.77   | 25.47     | 5,995,366  | 3,299,399  |

| Executor | CPU Time | Wall Time | Bytes Sent | Bytes Recv |
|----------|----------|-----------|------------|------------|
| LP-2010  | 55.90    | 27.45     | 737,109    | 7,648,074  |
| L-2013   | 101.69   | 42.05     | 980,193    | 4,693,761  |
| L-HKE    | 132.51   | 25.83     | 3,299,399  | 5,995,366  |

L-2013 runs substantially slower than LP-2010 due to the relative cost of the sub-computation being high.

## 32-bit Addition Circuit (439 gates)

| Builder | CPU Time | Wall Time | Bytes Sent | Bytes Recv |
|---------|----------|-----------|------------|------------|
| LP-2010 | 113.96   | 27.41     | 7,648,074  | 737,109    |
| L2013   | 171.21   | 42.03     | 4,693,761  | 980,193    |
| HKE     | 45.59    | 6.77      | 3,143,383  | 3,143,366  |
| L-HKE   | 145.77   | 25.47     | 5,995,366  | 3,299,399  |

| Executor | CPU Time | Wall Time | Bytes Sent | Bytes Recv |
|----------|----------|-----------|------------|------------|
| LP-2010  | 55.90    | 27.45     | 737,109    | 7,648,074  |
| L-2013   | 101.69   | 42.05     | 980,193    | 4,693,761  |
| L-HKE    | 132.51   | 25.83     | 3,299,399  | 5,995,366  |

L-HKE's results are very hopeful compared to LP-2010 and L-2013.

## AES-128 Encryption Circuit (33,872 gates)

| Builder | CPU Time | Wall Time | Bytes Sent  | Bytes Recv  |
|---------|----------|-----------|-------------|-------------|
| LP-2010 | 480.82   | 114.98    | 668,935,684 | 2,798,517   |
| L-2013  | 399.27   | 119.25    | 210,537,538 | 1,609,692   |
| HKE     | 185.47   | 32.95     | 238,300,835 | 238,300,840 |
| L-HKE   | 417.84   | 78.22     | 214,725,419 | 7,868,176   |

| Executor | CPU Time | Wall Time | Bytes Sent | Bytes Recv  |
|----------|----------|-----------|------------|-------------|
| LP-2010  | 227.91   | 116.15    | 2,798,517  | 668,935,684 |
| L-2013   | 270.99   | 119.27    | 1,609,692  | 210,537,538 |
| L-HKE    | 363.46   | 80.49     | 7,868,176  | 214,725,419 |

HKE still has a commanding lead. While L-2013 closes the gap to LP-2010.



## AES-128 Encryption Circuit (33,872 gates)

| Builder | CPU Time | Wall Time | Bytes Sent  | Bytes Recv  |
|---------|----------|-----------|-------------|-------------|
| LP-2010 | 480.82   | 114.98    | 668,935,684 | 2,798,517   |
| L-2013  | 399.27   | 119.25    | 210,537,538 | 1,609,692   |
| HKE     | 185.47   | 32.95     | 238,300,835 | 238,300,840 |
| L-HKE   | 417.84   | 78.22     | 214,725,419 | 7,868,176   |

| Executor | CPU Time | Wall Time | Bytes Sent | Bytes Recv  |
|----------|----------|-----------|------------|-------------|
| LP-2010  | 227.91   | 116.15    | 2,798,517  | 668,935,684 |
| L-2013   | 270.99   | 119.27    | 1,609,692  | 210,537,538 |
| L-HKE    | 363.46   | 80.49     | 7,868,176  | 214,725,419 |

At first glance L-HKE is hopeful. However, the timing break downs shows cause for concern.

---

## Some Conclusions

- ✿ A clear victory for HKE in Wall and CPU time, despite needing to build/check more circuits. Symmetric Cut and Choose merits further work.
- ✿ L-2013 is the protocol of choice if bandwidth is at a premium, e.g, when on a metered connection.
- ✿ L-HKE shows promise on smaller circuits over L-2013, but is hampered on large circuits by the optimisation roll backs.
- ✿ The primary variable for predicting cost is the number of inputs.
- ✿ The depth of the circuit (size of circuit minus inputs) mainly affects bandwidth.

---

# Achievements

- ✿ We provide the first implementations of LP-2010 and L-2013.
- ✿ We propose and implement a variant of L-2013 using HKE for the sub-computation.
- ✿ We provide the first practical comparison of LP-2010, L-2013, HKE and our variant.
- ✿ We suggest abandoning the Zero Knowledge Proof approach to consistency proving in exchange for the logarithm approach.